



Let's talk orchestration.

An Intro to scheduling dbt with Dagster.

Sydney dbt meetup October 2021
Benoit Perigaud



Agenda

An intro to dagster and dbt

- 1 What are we going to talk about
- 2 Why we might need orchestration tools
- 3 Dagster, Prefect and Airflow
- 4 Dagster core concepts
- 5 Demo time!



Benoit Perigaud

Me, myself and I

- Principal Analytics Engineer at Aginic
- Data consulting company, part of the list of dbt Preferred Consulting Providers



What this talk will and won't be about

Especially won't



- An intro to why you might need an orchestration tool
- A brief overview of different options in the data space
- An overview of key concepts in Dagster
- A repo that you can clone and play with



- CI/CD orchestration of dbt
- An in-depth view of which orchestration tool is the best for you
- Ins and Outs of Dagster
- How to build a Dagster project using best practices
- How to deploy Dagster



A story about dbt scheduling

You will most likely need a tool at some point

- No longer possible to do a simple dbt run followed by a dbt test
 - Snapshots
 - Different timing requirements
 - Near real-time analytics with incremental models
- Need observability and alerts



Different options

But where to start

- dbt Cloud
- automation servers
 - As part of CI/CD
- cloud capabilities
 - Azure example: Azure Function App / Azure Data Factory
 - ...
- cron 🤖
- data orchestration tools
 - Airflow
 - Dagsster
 - Prefect



The 3 main players for data orchestration

Elevator pitch



"Airflow is a platform created by the community to programmatically author, schedule and monitor workflows."

- The OG
- Big community and numerous packages
- More difficult to host/manage (managed services exist)
- Local development not as easy
- Existing dbt examples



"Dagster is a data orchestrator for machine learning, analytics, and ETL"

- In the process of releasing a commercial hosted version, on top of the open source package
- Cares about the "what" on top of the "how"
- Configurable IO Managers
- Concept of data assets
- Easy local testing
- Integrates with Airflow



"Orchestrate the modern data stack"

- Commercial cloud offering as well as open source
- Talks to "negative/positive engineering"
- More modern, like Dagster


Dagster 101

Main concepts

The current way of working (as of oct 2021)

- solid
 - functional unit of work
 - read inputs, do an action, emit outputs
 - can materialise assets
- pipeline
 - links solids together
 - e.g.: `solid_2(solid_1())`
- preset/mode
 - run pipelines with specific configuration

This is changing




<https://dagster.io/blog/core-apis-moving-towards-1-0>

Other dagster concepts

Differentiators

IO Managers

Ability to switch between In Memory, Local, Blob or create your own. Allow partial re-runs



Asset Materialisations

"Asset" is Dagster's word for an entity, external to solids, that is mutated or created by a solid.

Accessible through the Asset Catalog

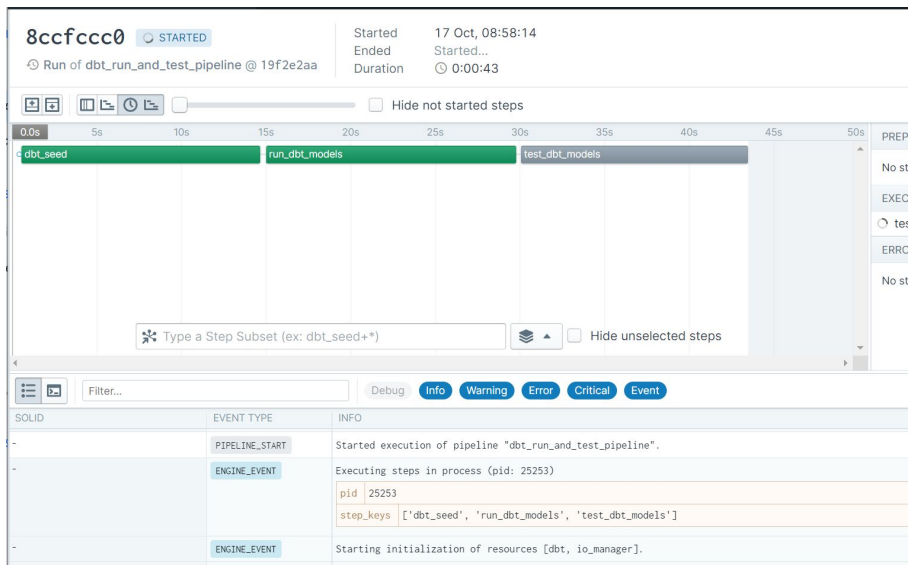


Demo!

dbt deps | dbt run | dbt test

- dagit shows the results and progress of the different dbt commands
- Ability to look at past runs and re-run them

```
> @pipeline(...  
)  
def dbt_run_and_test_pipeline():  
    dbt_seed_result = dbt_seed()  
    run_result = run_dbt_models(dbt_seed_result)  
    test_dbt_models(run_result)
```



Demo!

dbt assets materialisation

- Out of the box support for Asset Materialisation of dbt models
- Reads and Parses the “run_results.json” artefact
- By default, no support for tests, but very easy to implement (see function

generate_materializations dbt test)

```
@solid(  
    required_resource_keys={"dbt"},  
    config_schema={"models": Noneable(str), "exclude": Noneable(str)},  
)  
def run_dbt_models(context, run_result: DbtCliOutput) -> DbtCliOutput:  
    dbt_output = context.resources.dbt.run(  
        context.solid_config["models"], context.solid_config["exclude"]  
    )  
    for materialization in generate_materializations(dbt_output):  
        yield materialization  
    yield Output(dbt_output)
```

Materializations over Time

Graphs List

Run	16276c45	253f5291	d778a34	ceb267db	9b5cd522	f4064978	9ddf0e4d	8ccfcccc
Timestamp	16 Oct, 17:10	16 Oct, 17:56	16 Oct, 18:05	17 Oct, 07:55	17 Oct, 08:01	17 Oct, 08:26	17 Oct, 08:32	17 Oct, 08:58
Execution Time (s...)	0.01016235...	0.171459179980...	0.2045524120330...	0.1844639778137...	0.1665880680084...	0.1776683330535...	0.244520902633...	0.183441162109375
Compilation Start...	16T06:04...	2021-10-16T06:10...	2021-10-16T06:56...	2021-10-16T07:05...	2021-10-16T20:55...	2021-10-16T21:01...	2021-10-16T21:26...	2021-10-16T21:58...
Compilation Comp...	16T06:04...	2021-10-16T06:10...	2021-10-16T06:56...	2021-10-16T07:05...	2021-10-16T20:55...	2021-10-16T21:01...	2021-10-16T21:26...	2021-10-16T21:58...
Compilation Durati...	0.017297	0.017834	0.019963	0.017689	0.019061	0.034425	0.024666	0.019618
Execution Started ...	16T06:04...	2021-10-16T06:10...	2021-10-16T06:56...	2021-10-16T07:05...	2021-10-16T20:55...	2021-10-16T21:01...	2021-10-16T21:26...	2021-10-16T21:58...
Execution Comple...	16T06:04...	2021-10-16T06:10...	2021-10-16T06:56...	2021-10-16T07:05...	2021-10-16T20:55...	2021-10-16T21:01...	2021-10-16T21:26...	2021-10-16T21:58...
Step Execution Ti...	0:00:15	0:00:15	0:00:15	0:00:15	0:00:15	0:00:15	0:00:15	0:00:15

Compilation Duration



Execution Duration



Demo!

setup presets and schedules

- presets
 - used to execute ad-hoc pipelines
 - can be loaded as YAML files
- schedules
 - cron-based
 - able to "skip" runs based on filters

```
@pipeline(
    mode_defs=[ModeDefinition(resource_defs={"dbt": my_dbt_resource})],
    tags={"use-case": "dbt_run"},
    preset_defs=[
        PresetDefinition(
            "daily_schedule",
            run_config={
                "solids": {
                    "run_dbt_models": {"config": {"exclude": "tag:monthly+"}},
                    "test_dbt_models": {"config": {"exclude": "tag:monthly+"}},
                }
            },
        ),
        PresetDefinition(
            "monthly_schedule",
            run_config={
                "solids": {
                    "run_dbt_models": {"config": {"models": "tag:monthly+"}},
                    "test_dbt_models": {"config": {"models": "tag:monthly+"}},
                }
            },
        ),
    ],
)

def dbt_run_and_test_pipeline():
    dbt_seed_result = dbt_seed()
    run_result = run_dbt_models(dbt_seed_result)
    test_dbt_models(run_result)
```

```
@schedule(
    # 8am every day
    cron_schedule="0 8 * * *",
    pipeline_name="dbt_run_and_test_pipeline",
    execution_timezone="Australia/Sydney",
    should_execute=hour_filter,
)

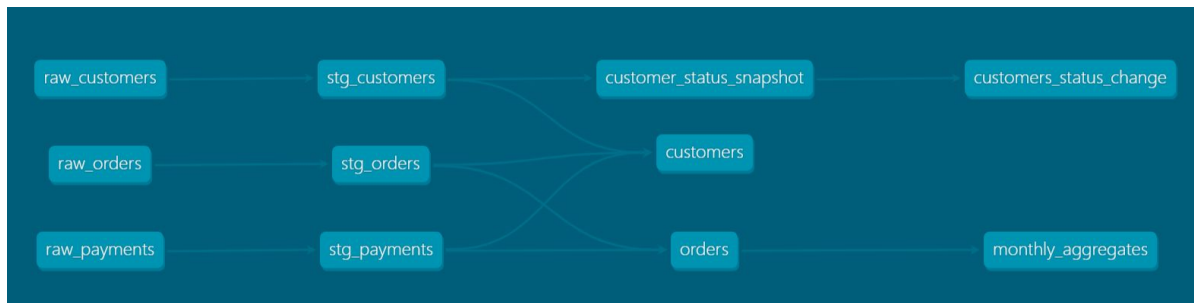
def dbt_daily_schedule(date):
    exclude_models = "tag:monthly+"
    return {
        "solids": {
            "run_dbt_models": {"config": {"exclude": exclude_models}},
            "test_dbt_models": {"config": {"exclude": exclude_models}},
        }
    }
```

Demo!

asset sensors

- trigger actions when an asset has been updated
- useful for dbt snapshots
- allows easier maintenance and less dependencies of cross-teams data pipelines

```
@asset_sensor(  
    asset_key=AssetKey(["model", "jaffle_shop", "stg_customers"]),  
    pipeline_name="dbt_snapshot_and_run_downstream_pipeline",  
)  
def asset_sensor_customer_snapshot(context, asset_event):  
    snapshot_name = "customer_status_snapshot"  
    yield RunRequest(  
        run_key=context.cursor,  
        run_config={  
            "solids": {  
                "dbt_snapshot": {"config": {"select": snapshot_name}},  
                "run_dbt_models": {"config": {"models": f"{snapshot_name}+"}},  
            }  
        },  
    )
```



What next?

More configuration and going beyond dbt



- Leverage “dbt build” 🎉
- Prevent concurrent runs of dbt
 - <https://docs.dagster.io/deployment/run-coordinator>
- Setup alerts and/or retries when the pipeline fails
- Maybe try to generate the full dbt DAG in dagster
 - Trying with Dynamic Mapping & Collect:
<https://docs.dagster.io/concepts/solids-pipelines/pipelines#dynamic-mapping--collect>
- Make our DAG more than dbt
 - Pause/resume EL pipeline
 - Refresh downstream applications (e.g. Power BI dataset)
 - Send notifications in slack
- Use the GraphQL API
 - Query information
 - Trigger pipelines



Thanks everyone!

That's all folks



- This talk (slides and code)
 - <https://github.com/Aginic/syd-dbt-meetup-dagster>
- Dagster
 - dagster: <https://dagster.io/>
 - dagster and dbt: <https://docs.dagster.io/integrations/dbt>
 - dagster slack: <https://dagster-slackin.herokuapp.com/>
- dbt
 - dbt build: <https://docs.getdbt.com/reference/commands/build>



Questions?

